

ex4_sol

May 8, 2019

1 Exercise 4 - Fully Connected Networks and the MNIST dataset

This exercise is based on <https://github.com/leriomaggio/deep-learning-keras-tensorflow>

2 The MNIST database

The MNIST (Modified National Institute of Standards and Technology) database ([link](#)) has a database of handwritten digits. The dataset consists of 28x28 grayscale images of the 10 digits.

Since this dataset is **provided** with Keras, we just ask the `keras.datasets` model for training and test data.

```
from keras.datasets import mnist          (X_train, y_train), (X_test, y_test) =  
mnist.load_data()
```

The training set has 60,000 samples. The test set has 10,000 samples. The digits are size-normalized and centered in a fixed-size image. The data page has a description on how the data was collected. It also reports the performance of various algorithms on the test dataset.

2.1 Task 1: Data preparation

- Download the data
- Inspect the data and plot a few of the images using `matplotlib.pyplot.imshow`
- Reshape the input data to be in vectorial form (original data are images)
- Convert the input data to do dtype `float32` using `astype` in order to scale it afterwards
- Normalize the design matrix to values between 0 and 1.
- How many classes do you have? How much data of each class?
- Convert the class vector to binary class matrices (**one-hot-vector**). Use the `to_categorical` function from `keras.utils` to convert integer labels to **one-hot-vectors**.
- Split the training set into training and validation data (30%)

2.1.1 Download the datasets and display first entry

```
In [1]: from keras.datasets import mnist
```

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Using TensorFlow backend.

```
In [2]: X_train[1]
```

```

Out[2]: array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0, 51, 159, 253, 159, 50,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0, 48, 238, 252, 252, 252, 237,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  54, 227, 253, 252, 239, 233, 252, 57,  6,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 10, 60,
                  224, 252, 253, 252, 202, 84, 252, 253, 122,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 163, 252,
                  252, 252, 253, 252, 252, 96, 189, 253, 167,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 51, 238, 253,
                  253, 190, 114, 253, 228, 47, 79, 255, 168,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0, 48, 238, 252, 252,
                  179, 12, 75, 121, 21,  0,  0, 253, 243, 50,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0, 38, 165, 253, 233, 208,
                  84,  0,  0,  0,  0,  0,  0, 253, 252, 165,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  7, 178, 252, 240, 71, 19,
                  28,  0,  0,  0,  0,  0,  0, 253, 252, 195,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0, 57, 252, 252, 63,  0,  0,
                  0,  0,  0,  0,  0,  0, 253, 252, 195,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0, 198, 253, 190,  0,  0,  0,
                  0,  0,  0,  0,  0,  0, 255, 253, 196,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0, 76, 246, 252, 112,  0,  0,  0,
                  0,  0,  0,  0,  0,  0, 253, 252, 148,  0,  0,  0,
                  0,  0],

```

```

[ 0, 0, 0, 0, 0, 0, 85, 252, 230, 25, 0, 0, 0,
 0, 0, 0, 0, 0, 7, 135, 253, 186, 12, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 85, 252, 223, 0, 0, 0, 0,
 0, 0, 0, 0, 7, 131, 252, 225, 71, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 85, 252, 145, 0, 0, 0, 0,
 0, 0, 0, 48, 165, 252, 173, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 86, 253, 225, 0, 0, 0, 0,
 0, 0, 114, 238, 253, 162, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 85, 252, 249, 146, 48, 29, 85,
 178, 225, 253, 223, 167, 56, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 85, 252, 252, 252, 229, 215, 252,
 252, 252, 196, 130, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 28, 199, 252, 252, 253, 252, 252,
 233, 145, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 25, 128, 252, 253, 252, 141,
 37, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0]], dtype=uint8)


```

```

In [3]: %matplotlib inline
import matplotlib.pyplot as plt
import matplotlib.cm as cm

fig = plt.figure()
plt.imshow(X_train[1], cmap=cm.Greys)
plt.show()

```



```
In [4]: plt.figure(figsize=(12,10))
        x, y = 10, 4
        for i in range(40):
            plt.subplot(y, x, i+1)
            plt.imshow(X_train[i], interpolation='nearest', cmap=cm.Greys)
        plt.show()
```



2.1.2 Reshape the datasets

In [5]: `X_train.shape`

Out[5]: (60000, 28, 28)

In [6]: `X_train = X_train.reshape(60000, 784)`
`X_test = X_test.reshape(10000, 784)`

In [7]: `X_train.shape`

Out[7]: (60000, 784)

2.1.3 Convert to float32

In [8]: `X_train.dtype`

Out[8]: dtype('uint8')

In [9]: `X_train = X_train.astype("float32")`
`X_test = X_test.astype("float32")`

In [10]: `X_train.dtype`

Out[10]: dtype('float32')

```
In [11]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler(feature_range=(0, 1))

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

[illegible]

0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.03921569, 0.23529413, 0.87843144,
 0.98823535, 0.9921569 , 0.98823535, 0.79215693, 0.32941177,
 0.98823535, 0.9921569 , 0.4784314 , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0.6392157 , 0.98823535, 0.98823535, 0.98823535, 0.9921569 ,
 0.98823535, 0.98823535, 0.37647063, 0.7411765 , 0.9921569 ,
 0.654902 , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.20000002, 0.9333334 , 0.9921569 ,
 0.9921569 , 0.74509805, 0.44705886, 0.9921569 , 0.8941177 ,
 0.18431373, 0.30980393, 1. , 0.65882355, 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.18823531,
 0.9333334 , 0.98823535, 0.98823535, 0.7019608 , 0.04705883,
 0.29411766, 0.47450984, 0.08235294, 0. , 0. ,
 0.9921569 , 0.95294124, 0.19607845, 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.14901961, 0.64705884, 0.9921569 , 0.91372555,
 0.81568635, 0.32941177, 0. , 0. , 0. ,
 0. , 0. , 0. , 0.9921569 , 0.98823535,
 0.64705884, 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0.02745098, 0.69803923,
 0.98823535, 0.94117653, 0.2784314 , 0.07450981, 0.10980393,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.9921569 , 0.98823535, 0.76470596, 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.22352943, 0.98823535, 0.98823535, 0.24705884,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.9921569 ,
 0.98823535, 0.76470596, 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.77647066,
 0.9921569 , 0.74509805, 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 1. , 0.9921569 , 0.7686275 ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.29803923, 0.96470594, 0.98823535, 0.43921572,
 0. , 0. , 0. , 0. , 0. ,

0. , 0. , 0. , 0. , 0. ,
 0.9921569 , 0.98823535, 0.5803922 , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.33333334,
 0.98823535, 0.90196085, 0.09803922, 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.02745098, 0.5294118 , 0.9921569 , 0.7294118 ,
 0.04705883, 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.33333334, 0.98823535, 0.8745099 ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0.02745098, 0.5137255 ,
 0.98823535, 0.882353 , 0.2784314 , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0.33333334, 0.98823535, 0.5686275 , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0.18823531, 0.64705884, 0.98823535, 0.6784314 , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0.3372549 , 0.9921569 ,
 0.882353 , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.44705886, 0.9333334 , 0.9921569 ,
 0.63529414, 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.33333334, 0.98823535, 0.97647065, 0.57254905,
 0.18823531, 0.1137255 , 0.33333334, 0.69803923, 0.882353 ,
 0.9921569 , 0.8745099 , 0.654902 , 0.21960786, 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0.33333334,
 0.98823535, 0.98823535, 0.98823535, 0.8980393 , 0.8431373 ,
 0.98823535, 0.98823535, 0.98823535, 0.7686275 , 0.50980395,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0.10980393, 0.7803922 , 0.98823535,
 0.98823535, 0.9921569 , 0.98823535, 0.98823535, 0.91372555,
 0.5686275 , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0.09803922, 0.5019608 , 0.98823535, 0.9921569 ,
 0.98823535, 0.5529412 , 0.14509805, 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,
 0. , 0. , 0. , 0. , 0. ,

[illegible]


2.1.5 Convert class vectors to binary class matrices

```
In [13]: y_train
```

```
Out[13]: array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

```
In [14]: import seaborn as sns
sns.countplot(y_train)
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2eee6c6c90>
```



```
In [15]: from keras.utils import np_utils
```

```
Y_train = np_utils.to_categorical(y_train, 10)  
Y_test = np_utils.to_categorical(y_test, 10)
```

```
In [16]: Y_train[1]
```

```
Out[16]: array([1., 0., 0., 0., 0., 0., 0., 0., 0., 0.], dtype=float32)
```


Split Training and Validation Data

```
In [17]: from sklearn.model_selection import train_test_split
```

```
X_train, X_val, Y_train, Y_val = train_test_split(X_train, Y_train, test_size=0.3, ra
```

```
In [18]: plt.imshow(X_train[1].reshape(28, 28), cmap=cm.Greys)
```

```
Out[18]: <matplotlib.image.AxesImage at 0x7f2eee634390>
```




```
In [19]: import numpy as np
         print(np.asarray(range(10)))
         print(Y_train[1].astype('int'))
```

```
[0 1 2 3 4 5 6 7 8 9]
[1 0 0 0 0 0 0 0 0 0]
```

```
In [20]: plt.imshow(X_val[0].reshape(28, 28), cmap=cm.Greys)
```

```
Out[20]: <matplotlib.image.AxesImage at 0x7f2eee4e3a10>
```



```
In [21]: print(np.asarray(range(10)))
         print(Y_val[0].astype('int'))
```

```
[0 1 2 3 4 5 6 7 8 9]
[0 0 0 0 1 0 0 0 0 0]
```

2.2 Task 2: Build and train a neural network

- Design a dense neural network structure.
- Choose softmax as activation for the output node (normalized multi-class probability)
- Use categorical_crossentropy as loss function (multi-class version of crossentropy)
- Use adam as optimizer and a batch size of 512 (speed things up)
- Train the NN over 50 epochs and plot the evolution of the training and validation loss as well as of one meaningful metric. What do you observe?
- Evaluate the performance on the test set using sklearn.metrics
- Plot the probability of being a *Zero* for true zeros and for all other numbers

2.3 Training

```
In [22]: from keras.models import Sequential
         from keras.layers.core import Dense

         model = Sequential()
         model.add(Dense(512, activation='relu', input_dim=784))
         model.add(Dense(256, activation='relu'))
```

```

model.add(Dense(10, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

model.summary()

```

WARNING:tensorflow:From /home/nackenho/miniconda/envs/TUDortmundMLSeminar/lib/python2.7/site-p
Instructions for updating:
Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 256)	131328
dense_3 (Dense)	(None, 10)	2570

Total params: 535,818
Trainable params: 535,818
Non-trainable params: 0

```
In [23]: history = model.fit(X_train, Y_train, batch_size=512, epochs=50, verbose=1, validation
```

WARNING:tensorflow:From /home/nackenho/miniconda/envs/TUDortmundMLSeminar/lib/python2.7/site-p
Instructions for updating:

Use tf.cast instead.

Train on 42000 samples, validate on 18000 samples

Epoch 1/50

42000/42000 [=====] - 2s 52us/step - loss: 0.4474 - acc: 0.8760 - val.

Epoch 2/50

42000/42000 [=====] - 2s 38us/step - loss: 0.1640 - acc: 0.9515 - val.

Epoch 3/50

42000/42000 [=====] - 2s 39us/step - loss: 0.1088 - acc: 0.9695 - val.

Epoch 4/50

42000/42000 [=====] - 2s 38us/step - loss: 0.0784 - acc: 0.9772 - val.

Epoch 5/50

42000/42000 [=====] - 2s 41us/step - loss: 0.0619 - acc: 0.9815 - val.

Epoch 6/50

42000/42000 [=====] - 2s 42us/step - loss: 0.0408 - acc: 0.9885 - val.

Epoch 7/50

42000/42000 [=====] - 2s 47us/step - loss: 0.0385 - acc: 0.9890 - val.

Epoch 8/50

42000/42000 [=====] - 2s 50us/step - loss: 0.0237 - acc: 0.9940 - val.

Epoch 9/50

42000/42000 [=====] - 2s 50us/step - loss: 0.0168 - acc: 0.9962 - val.

```

Epoch 10/50
42000/42000 [=====] - 2s 47us/step - loss: 0.0320 - acc: 0.9913 - val.
Epoch 11/50
42000/42000 [=====] - 2s 47us/step - loss: 0.0108 - acc: 0.9978 - val.
Epoch 12/50
42000/42000 [=====] - 3s 62us/step - loss: 0.0112 - acc: 0.9975 - val.
Epoch 13/50
42000/42000 [=====] - 2s 56us/step - loss: 0.0050 - acc: 0.9995 - val.
Epoch 14/50
42000/42000 [=====] - 2s 53us/step - loss: 0.0040 - acc: 0.9996 - val.
Epoch 15/50
42000/42000 [=====] - 2s 59us/step - loss: 0.0028 - acc: 0.9998 - val.
Epoch 16/50
42000/42000 [=====] - 3s 62us/step - loss: 0.0020 - acc: 0.9999 - val.
Epoch 17/50
42000/42000 [=====] - 2s 55us/step - loss: 0.0017 - acc: 1.0000 - val.
Epoch 18/50
42000/42000 [=====] - 2s 48us/step - loss: 0.0013 - acc: 1.0000 - val.
Epoch 19/50
42000/42000 [=====] - 2s 46us/step - loss: 0.0011 - acc: 1.0000 - val.
Epoch 20/50
42000/42000 [=====] - 2s 40us/step - loss: 9.3108e-04 - acc: 1.0000 -
Epoch 21/50
42000/42000 [=====] - 2s 37us/step - loss: 8.9243e-04 - acc: 1.0000 -
Epoch 22/50
42000/42000 [=====] - 2s 38us/step - loss: 6.6004e-04 - acc: 1.0000 -
Epoch 23/50
42000/42000 [=====] - 2s 41us/step - loss: 6.1369e-04 - acc: 1.0000 -
Epoch 24/50
42000/42000 [=====] - 2s 39us/step - loss: 6.1002e-04 - acc: 1.0000 -
Epoch 25/50
42000/42000 [=====] - 2s 38us/step - loss: 4.5800e-04 - acc: 1.0000 -
Epoch 26/50
42000/42000 [=====] - 2s 41us/step - loss: 4.5487e-04 - acc: 1.0000 -
Epoch 27/50
42000/42000 [=====] - 2s 39us/step - loss: 3.6766e-04 - acc: 1.0000 -
Epoch 28/50
42000/42000 [=====] - 2s 39us/step - loss: 3.3032e-04 - acc: 1.0000 -
Epoch 29/50
42000/42000 [=====] - 2s 41us/step - loss: 2.9913e-04 - acc: 1.0000 -
Epoch 30/50
42000/42000 [=====] - 2s 48us/step - loss: 2.6924e-04 - acc: 1.0000 -
Epoch 31/50
42000/42000 [=====] - 3s 76us/step - loss: 3.8782e-04 - acc: 1.0000 -
Epoch 32/50
42000/42000 [=====] - 2s 51us/step - loss: 2.3664e-04 - acc: 1.0000 -
Epoch 33/50
42000/42000 [=====] - 2s 40us/step - loss: 2.2486e-04 - acc: 1.0000 -

```

```

Epoch 34/50
42000/42000 [=====] - 2s 41us/step - loss: 2.0623e-04 - acc: 1.0000 -
Epoch 35/50
42000/42000 [=====] - 2s 40us/step - loss: 1.7491e-04 - acc: 1.0000 -
Epoch 36/50
42000/42000 [=====] - 2s 40us/step - loss: 1.6008e-04 - acc: 1.0000 -
Epoch 37/50
42000/42000 [=====] - 2s 41us/step - loss: 1.4587e-04 - acc: 1.0000 -
Epoch 38/50
42000/42000 [=====] - 2s 39us/step - loss: 1.3546e-04 - acc: 1.0000 -
Epoch 39/50
42000/42000 [=====] - 2s 41us/step - loss: 1.3024e-04 - acc: 1.0000 -
Epoch 40/50
42000/42000 [=====] - 2s 42us/step - loss: 1.1657e-04 - acc: 1.0000 -
Epoch 41/50
42000/42000 [=====] - 2s 40us/step - loss: 1.0743e-04 - acc: 1.0000 -
Epoch 42/50
42000/42000 [=====] - 2s 43us/step - loss: 1.0068e-04 - acc: 1.0000 -
Epoch 43/50
42000/42000 [=====] - 2s 40us/step - loss: 9.4129e-05 - acc: 1.0000 -
Epoch 44/50
42000/42000 [=====] - 2s 40us/step - loss: 8.7896e-05 - acc: 1.0000 -
Epoch 45/50
42000/42000 [=====] - 2s 42us/step - loss: 8.1163e-05 - acc: 1.0000 -
Epoch 46/50
42000/42000 [=====] - 2s 40us/step - loss: 7.6318e-05 - acc: 1.0000 -
Epoch 47/50
42000/42000 [=====] - 2s 41us/step - loss: 7.1498e-05 - acc: 1.0000 -
Epoch 48/50
42000/42000 [=====] - 2s 42us/step - loss: 6.6597e-05 - acc: 1.0000 -
Epoch 49/50
42000/42000 [=====] - 2s 39us/step - loss: 6.2577e-05 - acc: 1.0000 -
Epoch 50/50
42000/42000 [=====] - 2s 41us/step - loss: 5.8324e-05 - acc: 1.0000 -

```

2.3.1 Plotting the network history

As seen before, the return value of the fit function is a `keras.callbacks.History` object which contains the entire history of training/validation loss and defined metric (accuracy) for each epoch. Let's define a function to plot the history:

```


In [24]: def plot_history(network_history):
          plt.figure()
          plt.xlabel('Epochs')
          plt.ylabel('Loss')
          plt.plot(network_history.history['loss'])
          plt.plot(network_history.history['val_loss'])


```

```
plt.legend(['Training', 'Validation'])

plt.figure()
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.plot(network_history.history['acc'])
plt.plot(network_history.history['val_acc'])
plt.legend(['Training', 'Validation'], loc='lower right')
plt.show()
```

In [25]: plot_history(history)





2.4 Evaluation

```
In [26]: loss_and_metrics = model.evaluate(X_test, Y_test, batch_size=256)
         print loss_and_metrics
```

```
10000/10000 [=====] - 0s 20us/step
[0.09614004611876453, 0.9797]
```

```
In [27]: # Predict the values from the test dataset
         Y_pred = model.predict(X_test)
         # Convert predictions classes to one hot vectors
         Y_cls = np.argmax(Y_pred, axis = 1)
         # Convert validation observations to one hot vectors
         Y_true = np.argmax(Y_test, axis = 1)
```

```
In [28]: from sklearn.metrics import classification_report
         print 'Classification Report:\n', classification_report(Y_true,Y_cls)
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.98	0.98	1032

3	0.98	0.98	0.98	1010
4	0.98	0.98	0.98	982
5	0.98	0.98	0.98	892
6	0.98	0.98	0.98	958
7	0.98	0.97	0.98	1028
8	0.97	0.97	0.97	974
9	0.98	0.98	0.98	1009
micro avg	0.98	0.98	0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

2.4.1 Plotting the normalized probability $\text{prob}_i = p_i / \text{sum}(p_i)$

In multi-class problems the interpretation of the n-dimensional output is not always trivial, in particular if an output activation function is used which can not be interpreted as a probability. If one is only interested in distinguishing two of the classes one could build the ratio of these two class responses in order to get the best discrimination. Similarly, one could weight different classes according their importance for the specific problem. Because we have used a softmax activation together with the categorical cross-entropy we can directly interpret our output as probabilities and don't need to normalize it. If that is not the case you can define a multi-class probability for instance in the following way:


```
In [29]: def prob_multiclass(Y_pred, index):
          n_cls = len(Y_pred[0])

          Y_prob=[]
          for i in range(len(Y_pred)):
              numerator=Y_pred[i,index]
              denominator=0.0
              for idx in range(n_cls):
                  denominator+=Y_pred[i,idx]

              Y_prob.append(numerator/denominator)

          return np.asarray(Y_prob)

In [30]: label=0
          Y_pred_prob = prob_multiclass(Y_pred, label)
          plt.hist(Y_pred_prob[Y_true == label], alpha=0.5, color='red', bins=10, log = True)
          plt.hist(Y_pred_prob[Y_true != label], alpha=0.5, color='blue', bins=10, log = True)
          plt.legend(['digit == 0', 'digit != 0'], loc='upper right')
          plt.xlabel('Probability of being 0')
          plt.ylabel('Number of entries')
          plt.show()
```



2.4.2 Plot the confusion matrix

A good way to show the performance of a multi-class output is the confusion matrix:
http://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html

In [31]: *#Note, this code is taken straight from the SKLEARN website, an nice way of viewing c*

```
import itertools
def plot_confusion_matrix(cm, classes,
                           normalize=True,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
```

```

# Loop over data dimensions and create text annotations.
fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")


plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')


```

```

In [32]: from sklearn.metrics import confusion_matrix
         # compute the confusion matrix
         confusion_mtx = confusion_matrix(Y_true, Y_cls)
         # plot the confusion matrix
         plt.figure(figsize=(8,8))
         plot_confusion_matrix(confusion_mtx, classes = range(10))
         plt.figure(figsize=(8,8))
         plot_confusion_matrix(confusion_mtx, classes = range(10), normalize=False)

```





2.4.3 Plot wrong associations

Errors are difference between predicted labels and true labels

```
In [33]: errors = (Y_cls - Y_true != 0)
```

```
Y_cls_errors = Y_cls[errors]
Y_pred_errors = Y_pred[errors]
Y_true_errors = Y_true[errors]
X_test_errors = X_test[errors]
```

Define plotting function

```
In [34]: def display_errors(errors_index,img_errors,pred_errors, obs_errors):
        """ This function shows 6 images with their predicted and real labels"""
        n = 0
        nrows = 2
        ncols = 3
        fig, ax = plt.subplots(nrows,ncols,sharex=True,sharey=True)
        for row in range(nrows):
            for col in range(ncols):
                error = errors_index[n]
                ax[row,col].imshow((img_errors[error]).reshape((28,28)), cmap=cm.Greys, i
                ax[row,col].set_title("Predicted label :{}\nTrue label :{}".format(pred_e
                n += 1
```

Rank errors by difference in probability

```
In [35]: # Probabilities of the wrong predicted numbers
        Y_pred_errors_prob = np.max(Y_pred_errors,axis = 1)


        # Predicted probabilities of the true values in the error set
        true_prob_errors = np.diagonal(np.take(Y_pred_errors, Y_true_errors, axis=1))

        # Difference between the probability of the predicted label and the true label
        delta_pred_true_errors = Y_pred_errors_prob - true_prob_errors

        # Sorted list of the delta prob errors
        sorted_dela_errors = np.argsort(delta_pred_true_errors)

        # Top 6 errors
        most_important_errors = sorted_dela_errors[-6:]

In [36]: # Show the top 6 errors
        display_errors(most_important_errors, X_test_errors, Y_cls_errors, Y_true_errors)
```



2.5 Using Dropout Layers

As we have learned last time, the trainings and validation loss of the fit history is not comparable when using dropout. We can define our own callback function which calculates the loss and metric after each epoch for any dataset

```
In [37]: from keras.callbacks import Callback

class HistoryEpoch(Callback):
    def __init__(self, data):
        self.data = data

    def on_train_begin(self, logs={}):
        self.loss = []
        self.acc = []

    def on_epoch_end(self, epoch, logs={}):
        x, y = self.data
        l, a = self.model.evaluate(x, y, verbose=0)
        self.loss.append(l)
        self.acc.append(a)
```


2.6 Task 3: Using regularizer

- Modify your previous example network by adding a Dropout layer after each hidden layer
- Add l2 regularization to the hidden layers
- Use the new defined HistoryEpoch for training, validation and test data set in order to save a comparable loss function and metric. This is done by e.g.: `train_hist=HistoryEpoch((X_train, Y_train))`. In the fit function you can call the callback then by specifying `callbacks=[train_hist]`.
- Plot the loss and metric evolution and compare the calculated loss with the default loss from the history
- Evaluate the performance of the NN as for the unregularized NN and compare the performance

```
In [38]: from keras.layers.core import Dropout
         from keras.regularizers import l2

         dropout=0.5
         l2_lambda = 0.0001

         model_dropout = Sequential()
         model_dropout.add(Dense(512, activation='relu', kernel_regularizer=l2(l2_lambda), input_shape=(X_train.shape[1],)))
         model_dropout.add(Dropout(dropout))
         model_dropout.add(Dense(256, activation='relu', kernel_regularizer=l2(l2_lambda)))
         model_dropout.add(Dropout(dropout))
         model_dropout.add(Dense(10, activation='softmax'))

         train_hist=HistoryEpoch((X_train, Y_train))
         val_hist=HistoryEpoch((X_val, Y_val))
         test_hist=HistoryEpoch((X_test, Y_test))

         model_dropout.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
         model_dropout.summary()
```

WARNING:tensorflow:From /home/nackenho/miniconda/envs/TUDortmundMLSeminar/lib/python2.7/site-packages/tensorflow/core/framework/op_def_registry.py:100: *tf.nn.dropout* is deprecated and will be removed in a future version.

Instructions for updating:

Please use ``rate`` instead of ``keep_prob``. Rate should be set to ``rate = 1 - keep_prob``.

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 256)	131328
dropout_2 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570

```

=====
Total params: 535,818
Trainable params: 535,818
Non-trainable params: 0
-----

```

```
In [39]: history_dropout = model_dropout.fit(X_train, Y_train, batch_size=512, epochs=50, verbose=0)
```

Train on 42000 samples, validate on 18000 samples

```

Epoch 1/50
42000/42000 [=====] - 3s 68us/step - loss: 0.7581 - acc: 0.7856 - val_loss: 0.7581
Epoch 2/50
42000/42000 [=====] - 2s 47us/step - loss: 0.3438 - acc: 0.9214 - val_loss: 0.3438
Epoch 3/50
42000/42000 [=====] - 2s 47us/step - loss: 0.2781 - acc: 0.9409 - val_loss: 0.2781
Epoch 4/50
42000/42000 [=====] - 2s 47us/step - loss: 0.2373 - acc: 0.9528 - val_loss: 0.2373
Epoch 5/50
42000/42000 [=====] - 2s 50us/step - loss: 0.2167 - acc: 0.9585 - val_loss: 0.2167
Epoch 6/50
42000/42000 [=====] - 2s 49us/step - loss: 0.2022 - acc: 0.9622 - val_loss: 0.2022
Epoch 7/50
42000/42000 [=====] - 2s 49us/step - loss: 0.1848 - acc: 0.9676 - val_loss: 0.1848
Epoch 8/50
42000/42000 [=====] - 2s 51us/step - loss: 0.1831 - acc: 0.9676 - val_loss: 0.1831
Epoch 9/50
42000/42000 [=====] - 2s 50us/step - loss: 0.1687 - acc: 0.9713 - val_loss: 0.1687
Epoch 10/50
42000/42000 [=====] - 2s 52us/step - loss: 0.1553 - acc: 0.9754 - val_loss: 0.1553
Epoch 11/50
42000/42000 [=====] - 2s 50us/step - loss: 0.1630 - acc: 0.9723 - val_loss: 0.1630
Epoch 12/50
42000/42000 [=====] - 2s 48us/step - loss: 0.1470 - acc: 0.9779 - val_loss: 0.1470
Epoch 13/50
42000/42000 [=====] - 2s 48us/step - loss: 0.1472 - acc: 0.9772 - val_loss: 0.1472
Epoch 14/50
42000/42000 [=====] - 2s 49us/step - loss: 0.1393 - acc: 0.9787 - val_loss: 0.1393
Epoch 15/50
42000/42000 [=====] - 2s 49us/step - loss: 0.1321 - acc: 0.9802 - val_loss: 0.1321
Epoch 16/50
42000/42000 [=====] - 2s 53us/step - loss: 0.1305 - acc: 0.9805 - val_loss: 0.1305
Epoch 17/50
42000/42000 [=====] - 2s 52us/step - loss: 0.1242 - acc: 0.9819 - val_loss: 0.1242
Epoch 18/50
42000/42000 [=====] - 2s 53us/step - loss: 0.1364 - acc: 0.9774 - val_loss: 0.1364
Epoch 19/50
42000/42000 [=====] - 3s 62us/step - loss: 0.1307 - acc: 0.9801 - val_loss: 0.1307

```

```

Epoch 20/50
42000/42000 [=====] - 3s 67us/step - loss: 0.1266 - acc: 0.9808 - val.
Epoch 21/50
42000/42000 [=====] - 3s 64us/step - loss: 0.1216 - acc: 0.9832 - val.
Epoch 22/50
42000/42000 [=====] - 2s 50us/step - loss: 0.1222 - acc: 0.9820 - val.
Epoch 23/50
42000/42000 [=====] - 2s 48us/step - loss: 0.1187 - acc: 0.9832 - val.
Epoch 24/50
42000/42000 [=====] - 2s 53us/step - loss: 0.1110 - acc: 0.9857 - val.
Epoch 25/50
42000/42000 [=====] - 2s 45us/step - loss: 0.1078 - acc: 0.9860 - val.
Epoch 26/50
42000/42000 [=====] - 2s 58us/step - loss: 0.1094 - acc: 0.9858 - val.
Epoch 27/50
42000/42000 [=====] - 2s 52us/step - loss: 0.1058 - acc: 0.9863 - val.
Epoch 28/50
42000/42000 [=====] - 2s 45us/step - loss: 0.1039 - acc: 0.9864 - val.
Epoch 29/50
42000/42000 [=====] - 2s 50us/step - loss: 0.1021 - acc: 0.9867 - val.
Epoch 30/50
42000/42000 [=====] - 2s 47us/step - loss: 0.1020 - acc: 0.9867 - val.
Epoch 31/50
42000/42000 [=====] - 2s 53us/step - loss: 0.1017 - acc: 0.9866 - val.
Epoch 32/50
42000/42000 [=====] - 3s 62us/step - loss: 0.0979 - acc: 0.9874 - val.
Epoch 33/50
42000/42000 [=====] - 2s 45us/step - loss: 0.1000 - acc: 0.9866 - val.
Epoch 34/50
42000/42000 [=====] - 2s 46us/step - loss: 0.1001 - acc: 0.9866 - val.
Epoch 35/50
42000/42000 [=====] - 2s 48us/step - loss: 0.1000 - acc: 0.9870 - val.
Epoch 36/50
42000/42000 [=====] - 2s 54us/step - loss: 0.0944 - acc: 0.9885 - val.
Epoch 37/50
42000/42000 [=====] - 2s 59us/step - loss: 0.1031 - acc: 0.9853 - val.
Epoch 38/50
42000/42000 [=====] - 2s 52us/step - loss: 0.0990 - acc: 0.9866 - val.
Epoch 39/50
42000/42000 [=====] - 3s 80us/step - loss: 0.0958 - acc: 0.9880 - val.
Epoch 40/50
42000/42000 [=====] - 3s 64us/step - loss: 0.0941 - acc: 0.9885 - val.
Epoch 41/50
42000/42000 [=====] - 2s 47us/step - loss: 0.0955 - acc: 0.9872 - val.
Epoch 42/50
42000/42000 [=====] - 2s 47us/step - loss: 0.1019 - acc: 0.9856 - val.
Epoch 43/50
42000/42000 [=====] - 2s 48us/step - loss: 0.0946 - acc: 0.9882 - val.

```

```

Epoch 44/50
42000/42000 [=====] - 2s 47us/step - loss: 0.1065 - acc: 0.9851 - val.
Epoch 45/50
42000/42000 [=====] - 2s 48us/step - loss: 0.0973 - acc: 0.9873 - val.
Epoch 46/50
42000/42000 [=====] - 2s 49us/step - loss: 0.0997 - acc: 0.9863 - val.
Epoch 47/50
42000/42000 [=====] - 2s 52us/step - loss: 0.0972 - acc: 0.9874 - val.
Epoch 48/50
42000/42000 [=====] - 2s 55us/step - loss: 0.1002 - acc: 0.9860 - val.
Epoch 49/50
42000/42000 [=====] - 2s 51us/step - loss: 0.0996 - acc: 0.9868 - val.
Epoch 50/50
42000/42000 [=====] - 2s 46us/step - loss: 0.0961 - acc: 0.9873 - val.



```

```



In [40]: plt.figure()
         plt.xlabel('Epochs')
         plt.ylabel('Loss')
         plt.plot(train_hist.loss)
         plt.plot(val_hist.loss)
         plt.plot(test_hist.loss)
         plt.legend(['Training', 'Validation', 'Testing'])

         plt.figure()
         plt.xlabel('Epochs')
         plt.ylabel('Accuracy')
         plt.plot(train_hist.acc)
         plt.plot(val_hist.acc)
         plt.plot(test_hist.acc)
         plt.legend(['Training', 'Validation', 'Testing'], loc='lower right')
         plt.show()

```



In [41]: `plot_history(history_dropout)`



2.7 Evaluation

```
In [42]: loss_and_metrics = model_dropout.evaluate(X_test, Y_test, batch_size=256)
        print loss_and_metrics
```

```
10000/10000 [=====] - 0s 21us/step
[0.1144553028523922, 0.9832]
```


```
In [43]: # Predict the values from the test dataset
        Y_pred = model_dropout.predict(X_test)
        # Convert predictions classes to one hot vectors
        Y_cls = np.argmax(Y_pred, axis = 1)
        # Convert validation observations to one hot vectors
        Y_true = np.argmax(Y_test, axis = 1)
```

```
In [44]: from sklearn.metrics import classification_report
        print 'Classification Report:\n', classification_report(Y_true,Y_cls)
```


Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.98	0.98	1032
3	0.97	0.99	0.98	1010
4	0.98	0.98	0.98	982
5	0.99	0.98	0.99	892
6	0.99	0.98	0.98	958
7	0.98	0.98	0.98	1028
8	0.98	0.98	0.98	974
9	0.98	0.98	0.98	1009
micro avg	0.98	0.98	0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

```
In [45]: label=0
        Y_pred_prob = prob_multiclass(Y_pred, label)
        plt.hist(Y_pred_prob[Y_true == label], alpha=0.5, color='red', bins=10, log = True)
        plt.hist(Y_pred_prob[Y_true != label], alpha=0.5, color='blue', bins=10, log = True)
        plt.legend(['digit == 0', 'digit != 0'], loc='upper right')
        plt.xlabel('Probability of being 0')
        plt.ylabel('Number of entries')
        plt.show()
```



```
In [46]: # compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_cls)
# plot the confusion matrix
plt.figure(figsize=(8,8))
plot_confusion_matrix(confusion_mtx, classes = range(10), normalize=False)
```

```
In [47]: #errors
errors = (Y_cls - Y_true != 0)

Y_cls_errors = Y_cls[errors]
Y_pred_errors = Y_pred[errors]
Y_true_errors = Y_true[errors]
X_test_errors = X_test[errors]

# Probabilities of the wrong predicted numbers
Y_pred_errors_prob = np.max(Y_pred_errors,axis = 1)

# Predicted probabilities of the true values in the error set
true_prob_errors = np.diagonal(np.take(Y_pred_errors, Y_true_errors, axis=1))
```


```

# Difference between the probability of the predicted label and the true label
delta_pred_true_errors = Y_pred_errors_prob - true_prob_errors

# Sorted list of the delta prob errors
sorted_delta_errors = np.argsort(delta_pred_true_errors)

# Top 6 errors
most_important_errors = sorted_delta_errors[-6:]
# Show the top 6 errors
display_errors(most_important_errors, X_test_errors, Y_cls_errors, Y_true_errors)

```



2.8 Early Stopping as a regularizer

- If you continue training, at some point the validation loss will start to increase: that is when the model starts to **overfit**. We can use EarlyStopping as a regularizer:

```
In [48]: from keras.callbacks import EarlyStopping
```

```

early_stop = EarlyStopping(monitor='val_loss', patience=5, verbose=1)
#Also possible choice:
#early_stop = EarlyStopping(monitor='val_acc', patience=5, verbose=1)

dropout=0.5

```

```

model_ES = Sequential()
model_ES.add(Dense(512, activation='relu', kernel_regularizer=l2(l2_lambda), input_dim=input_dim))
model_ES.add(Dropout(dropout))
model_ES.add(Dense(256, activation='relu', kernel_regularizer=l2(l2_lambda)))
model_ES.add(Dropout(dropout))
model_ES.add(Dense(10, activation='softmax'))

model_ES.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

model_dropout.summary()

```

```

-----
Layer (type)                 Output Shape              Param #
-----
dense_4 (Dense)              (None, 512)               401920
-----
dropout_1 (Dropout)          (None, 512)               0
-----
dense_5 (Dense)              (None, 256)              131328
-----
dropout_2 (Dropout)          (None, 256)               0
-----
dense_6 (Dense)              (None, 10)                2570
=====
Total params: 535,818
Trainable params: 535,818
Non-trainable params: 0
-----

```

```

In [49]: history_ES = model_ES.fit(X_train, Y_train, validation_data = (X_test, Y_test), epochs=100,
                                     callbacks=[early_stop])

```

Train on 42000 samples, validate on 10000 samples

```



Epoch 1/100
42000/42000 [=====] - 4s 87us/step - loss: 0.6242 - acc: 0.8320 - val_loss: 0.3106
Epoch 2/100
42000/42000 [=====] - 4s 87us/step - loss: 0.3106 - acc: 0.9324 - val_loss: 0.2539
Epoch 3/100
42000/42000 [=====] - 6s 143us/step - loss: 0.2539 - acc: 0.9476 - val_loss: 0.2207
Epoch 4/100
42000/42000 [=====] - 3s 80us/step - loss: 0.2207 - acc: 0.9571 - val_loss: 0.2044
Epoch 5/100
42000/42000 [=====] - 6s 152us/step - loss: 0.2044 - acc: 0.9634 - val_loss: 0.1868
Epoch 6/100
42000/42000 [=====] - 6s 153us/step - loss: 0.1868 - acc: 0.9670 - val_loss: 0.1868
Epoch 7/100

```

42000/42000 [=====] - 5s 130us/step - loss: 0.1763 - acc: 0.9695 - val.
Epoch 8/100
42000/42000 [=====] - 4s 107us/step - loss: 0.1702 - acc: 0.9708 - val.
Epoch 9/100
42000/42000 [=====] - 5s 109us/step - loss: 0.1663 - acc: 0.9726 - val.
Epoch 10/100
42000/42000 [=====] - 4s 85us/step - loss: 0.1586 - acc: 0.9736 - val.
Epoch 11/100
42000/42000 [=====] - 3s 77us/step - loss: 0.1509 - acc: 0.9762 - val.
Epoch 12/100
42000/42000 [=====] - 3s 83us/step - loss: 0.1475 - acc: 0.9766 - val.
Epoch 13/100
42000/42000 [=====] - 3s 77us/step - loss: 0.1439 - acc: 0.9783 - val.
Epoch 14/100
42000/42000 [=====] - 3s 72us/step - loss: 0.1460 - acc: 0.9770 - val.
Epoch 15/100
42000/42000 [=====] - 4s 90us/step - loss: 0.1365 - acc: 0.9800 - val.
Epoch 16/100
42000/42000 [=====] - 3s 80us/step - loss: 0.1343 - acc: 0.9803 - val.
Epoch 17/100
42000/42000 [=====] - 4s 88us/step - loss: 0.1408 - acc: 0.9780 - val.
Epoch 18/100
42000/42000 [=====] - 3s 70us/step - loss: 0.1368 - acc: 0.9795 - val.
Epoch 19/100
42000/42000 [=====] - 3s 69us/step - loss: 0.1296 - acc: 0.9813 - val.
Epoch 20/100
42000/42000 [=====] - 3s 74us/step - loss: 0.1253 - acc: 0.9815 - val.
Epoch 21/100
42000/42000 [=====] - 3s 83us/step - loss: 0.1262 - acc: 0.9827 - val.
Epoch 22/100
42000/42000 [=====] - 4s 84us/step - loss: 0.1235 - acc: 0.9831 - val.
Epoch 23/100
42000/42000 [=====] - 4s 88us/step - loss: 0.1257 - acc: 0.9815 - val.
Epoch 24/100
42000/42000 [=====] - 4s 100us/step - loss: 0.1262 - acc: 0.9815 - val.
Epoch 25/100
42000/42000 [=====] - 3s 83us/step - loss: 0.1230 - acc: 0.9829 - val.
Epoch 26/100
42000/42000 [=====] - 3s 70us/step - loss: 0.1217 - acc: 0.9832 - val.
Epoch 27/100
42000/42000 [=====] - 4s 83us/step - loss: 0.1200 - acc: 0.9838 - val.
Epoch 28/100
42000/42000 [=====] - 4s 91us/step - loss: 0.1244 - acc: 0.9823 - val.
Epoch 29/100
42000/42000 [=====] - 3s 71us/step - loss: 0.1235 - acc: 0.9829 - val.
Epoch 30/100
42000/42000 [=====] - 4s 95us/step - loss: 0.1245 - acc: 0.9821 - val.
Epoch 31/100

42000/42000 [=====] - 4s 85us/step - loss: 0.1201 - acc: 0.9839 - val.
Epoch 00031: early stopping

In [50]: plot_history(history_ES)



3 Bonus: Inspecting Layers

```
In [51]: # We already used `summary`  
         model_dropout.summary()
```

```
-----  
Layer (type)                 Output Shape                 Param #  
-----  
dense_4 (Dense)              (None, 512)                  401920  
-----  
dropout_1 (Dropout)          (None, 512)                  0  
-----  
dense_5 (Dense)              (None, 256)                  131328  
-----  
dropout_2 (Dropout)          (None, 256)                  0  
-----  
dense_6 (Dense)              (None, 10)                   2570  
-----  
Total params: 535,818  
Trainable params: 535,818  
Non-trainable params: 0  
-----
```

3.0.1 model.layers is iterable

```
In [52]: print('Model Input Tensors: ', model.input)  
         print('Layers - Network Configuration:')  
         for layer in model.layers:  
             print(layer.name, layer.trainable)  
             print('Layer Configuration:')  
             print(layer.get_config(), )  
         print('Model Output Tensors: ', model.output)  
  
(('Model Input Tensors: ', <tf.Tensor 'dense_1_input:0' shape=(?, 784) dtype=float32>)  
Layers - Network Configuration:  
(('dense_1', True)  
Layer Configuration:  
({'kernel_initializer': {'class_name': 'VarianceScaling', 'config': {'distribution': 'uniform'  
(('dense_2', True)  
Layer Configuration:  
({'kernel_initializer': {'class_name': 'VarianceScaling', 'config': {'distribution': 'uniform'  
(('dense_3', True)
```

Layer Configuration:

```
({'kernel_initializer': {'class_name': 'VarianceScaling', 'config': {'distribution': 'uniform'}  
(Model Output Tensors: ', <tf.Tensor 'dense_3/Softmax:0' shape=(?, 10) dtype=float32>)
```

3.1 Extract hidden layer representation of the given data

One **simple** way to do it is to use the weights of your model to build a new model that's truncated at the layer you want to read.

Then you can run the `._predict(X_batch)` method to get the activations for a batch of inputs.

```
In [53]: model_truncated = Sequential()  
         model_truncated.add(Dense(512, activation='relu', input_shape=(784,)))  
         model_truncated.add(Dropout(dropout))  
         model_truncated.add(Dense(256, activation='relu'))  
  
         for i, layer in enumerate(model_truncated.layers):  
             layer.set_weights(model_dropout.layers[i].get_weights())  
  
         model_truncated.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['a  
  
In [54]: # Check  
         np.all(model_truncated.layers[0].get_weights()[0] == model.layers[0].get_weights()[0])  
  
Out[54]: False  
  
In [55]: hidden_features = model_truncated.predict(X_train)  
  
In [56]: hidden_features.shape  
  
Out[56]: (42000, 256)  
  
In [57]: X_train.shape  
  
Out[57]: (42000, 784)
```

Hint: Alternative Method to get activations (Using `keras.backend` function on Tensors)

```
def get_activations(model, layer, X_batch):  
    activations_f = K.function([model.layers[0].input, K.learning_phase()], [layer.output,])  
    activations = activations_f((X_batch, False))  
    return activations
```

3.1.1 Generate the Embedding of Hidden Features

Dimensionality reduction to `dim=20` by using principal component analysis (PCA)

```
In [58]: from sklearn.decomposition import PCA  
         pca = PCA(n_components=20)  
         pca_result = pca.fit_transform(hidden_features)  
         print('Variance PCA: {}'.format(np.sum(pca.explained_variance_ratio_)))
```

Variance PCA: 0.95031619072

Dimensionality reduction to dim=2 by using t-distributed stochastic neighbor embedding (TSNE)

```
In [59]: from sklearn.manifold import TSNE
```

```
tsne = TSNE(n_components=2)
X_tsne = tsne.fit_transform(pca_result[:1000]) ## Reduced for computational issues
```

```
In [60]: colors_map = np.argmax(Y_train, axis=1)
```

```
In [61]: X_tsne.shape
```

```
Out[61]: (1000, 2)
```

```
In [62]: nb_classes=10
```

```
In [63]: np.where(colors_map==6)
```

```
Out[63]: (array([ 2, 8, 23, ..., 41927, 41983, 41988]),)
```

```
In [64]: colors = np.array([x for x in 'b-g-r-c-m-y-k-purple-coral-lime'.split('-')])
colors_map = np.argmax(Y_train, axis=1)
colors_map = colors_map[:1000]
plt.figure(figsize=(10,10))
for cl in range(nb_classes):
    indices = np.where(colors_map==cl)
    plt.scatter(X_tsne[indices,0], X_tsne[indices, 1], c=colors[cl], label=cl)
plt.legend()
plt.show()
```